

# Apache Spark Fundamentals

Dilip K. Prasad

# Outline

Introduction

Spark Fundamentals

Spark Architecture

Spark and it's Ecosystem

Spark vs Hadoop

RDD Fundamentals

Spark Transformations, Actions and Operations

# Introduction

2009 AMPLab (Berkley) – resource/cluster manager

2009                No YARN

MESOS

Created framework to test Mesos – Spark was created (which is a in memory execution) – huge RAM requirement

2010/11           SPARK further development started

2012               Apache Spark ->                               Databricks

2015               It got traction and more support system came up

2021               INF2220 Cloud and Big Data Technology started

Spark version 0

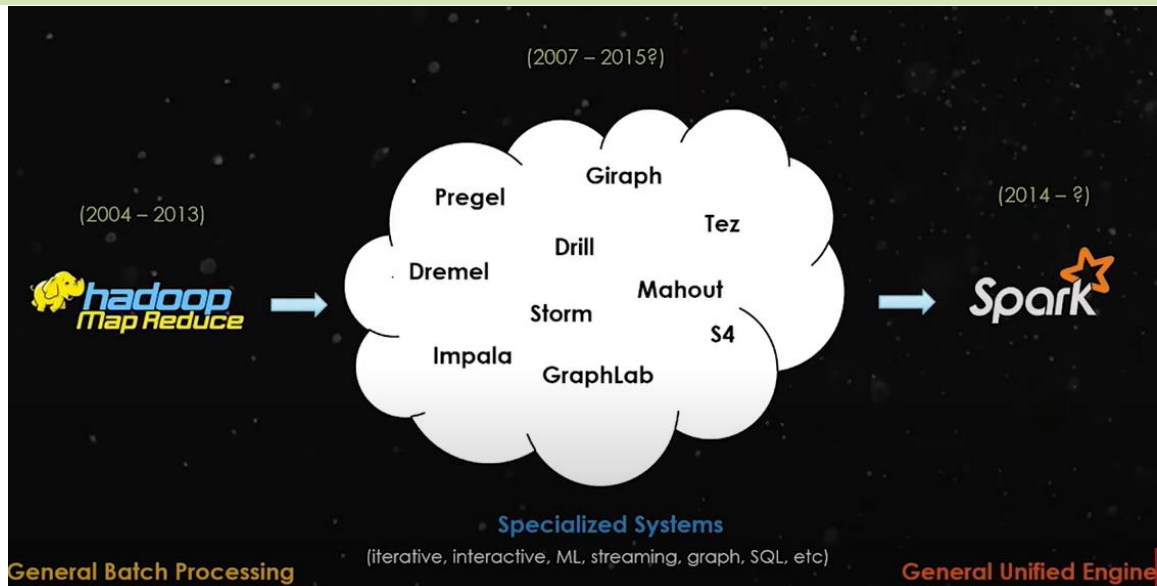
1.6

2.7

3.2.0

Best book – Spark the definitive guide – Mathe Zaharia (**don't buy** it its not text book type)

# Changing Big Data World



# Introduction to Spark



Scheduling




Monitoring

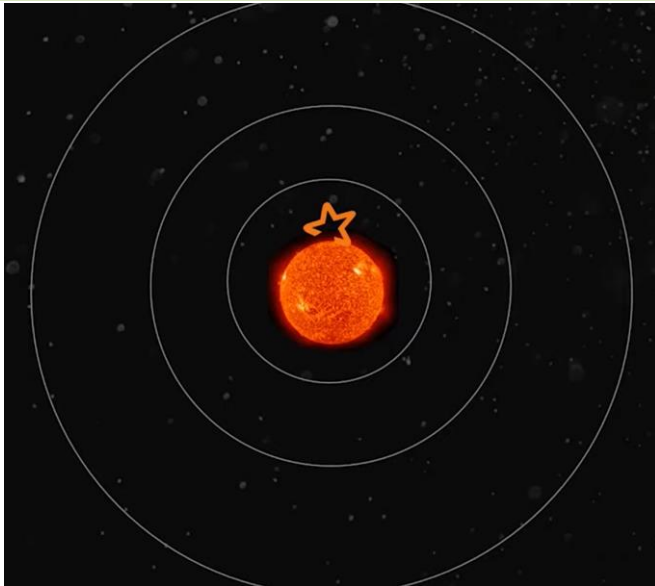


Distributing

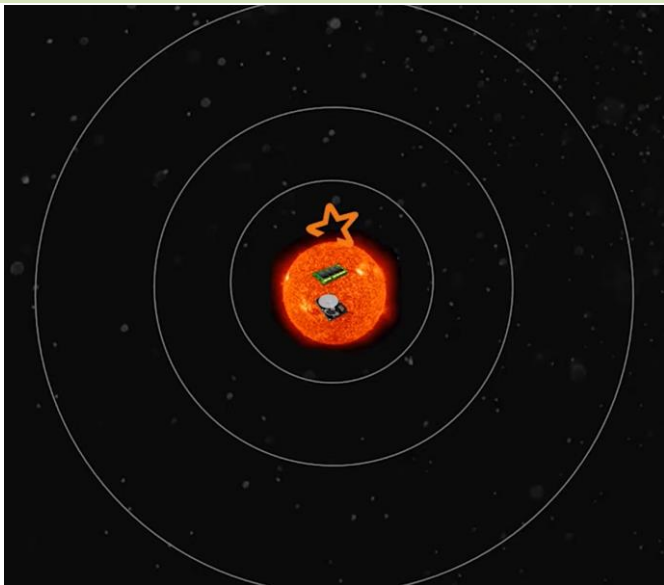
# Introduction to Spark



# Spark Architecture & its Ecosystem




# Spark Architecture & its Ecosystem





# Spark Architecture & its Ecosystem



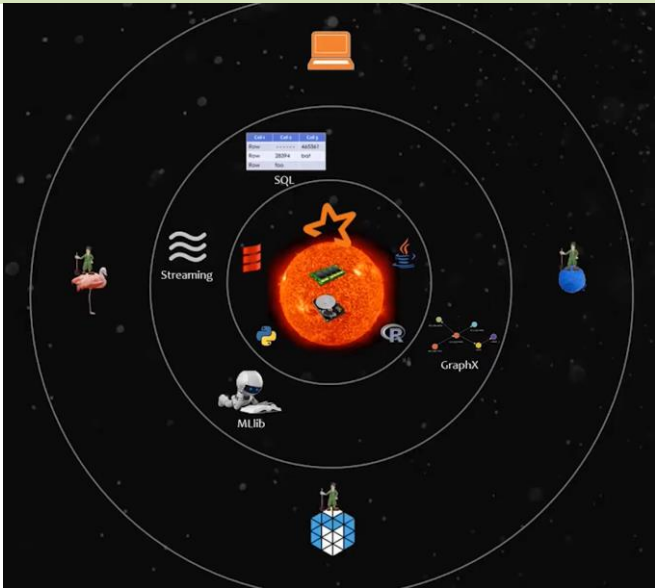
# Spark Architecture & its Ecosystem



# Spark Architecture & its Ecosystem



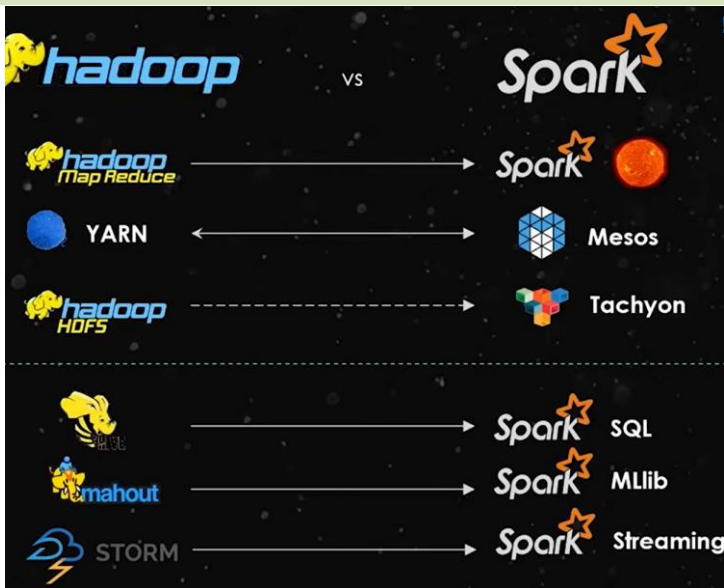
# Spark Architecture & its Ecosystem



# Spark Architecture & its Ecosystem



# Hadoop Vs Spark




## DISTRIBUTORS



## APPLICATIONS



# How spark program get executed





# Sort Competition

	Hadoop MR Record (2013)	Spark Record (2014)	Spark, 3x faster with 1/10 the nodes
Data Size	102.5 TB	100 TB	
Elapsed Time	72 mins	23 mins	
# Nodes	2100	206	
# Cores	50400 physical	6592 virtualized	
Cluster disk throughput	3150 GB/s (est.)	618 GB/s	
Network	dedicated data center, 10Gbps	virtualized (EC2) 10Gbps network	
<b>Sort rate</b>	<b>1.42 TB/min</b>	<b>4.27 TB/min</b>	
<b>Sort rate/node</b>	<b>0.67 GB/min</b>	<b>20.7 GB/min</b>	

Sort benchmark, Daytona Gray: sort of 100 TB of data (1 trillion records)

<http://databricks.com/blog/2014/11/05/spark-officially-sets-a-new-record-in-large-scale-sorting.html>

# Resilient Distributed Datasets (RDDs)

- RDDs (Resilient Distributed Datasets) is Data Containers
- All the different processing components in Spark share the same abstraction called RDD
- As applications share the RDD abstraction, you can mix different kind of transformations to create new RDDs
- Created by parallelizing a collection or reading a file
- Fault tolerant

# RDD



Error, ts, msg1  
Warn, ts, msg2  
Error, ts, msg1


Info, ts, msg8  
Warn, ts, msg2  
Info, ts, msg8

Error, ts, msg3  
Info, ts, msg5  
Info, ts, msg5


Error, ts, msg4  
Warn, ts, msg9  
Error, ts, msg1

logLinesRDD  
(input/base RDD)

# RDD



Data  
20GB



# DataFrames & SparkSQL

- DataFrames (DFs) is one of the other distributed datasets organized in named columns
- Similar to a relational database, Python Pandas Dataframe or R's DataTables
  - Immutable once constructed
  - Track lineage
  - Enable distributed computations
- How to construct Dataframes
  - Read from file(s)
  - Transforming an existing DFs(Spark or Pandas)
  - Parallelizing a python collection list
  - Apply transformations and actions

# DataFrame example

// Create a new DataFrame that contains “students”

```
students = users.filter(users.age < 21)
```

//Alternatively, using Pandas-like syntax

```
students = users[users.age < 21]
```

//Count the number of students users by gender


```
students.groupBy("gender").count()
```

// Join young students with another DataFrame called logs

```
students.join(logs, logs.userId == users.userId,  
"left_outer")
```

# RDDs vs. DataFrames

- RDDs provide a low level interface into Spark
- DataFrames have a schema
- DataFrames are cached and optimized by Spark
- DataFrames are built on top of the RDDs and the core Spark API

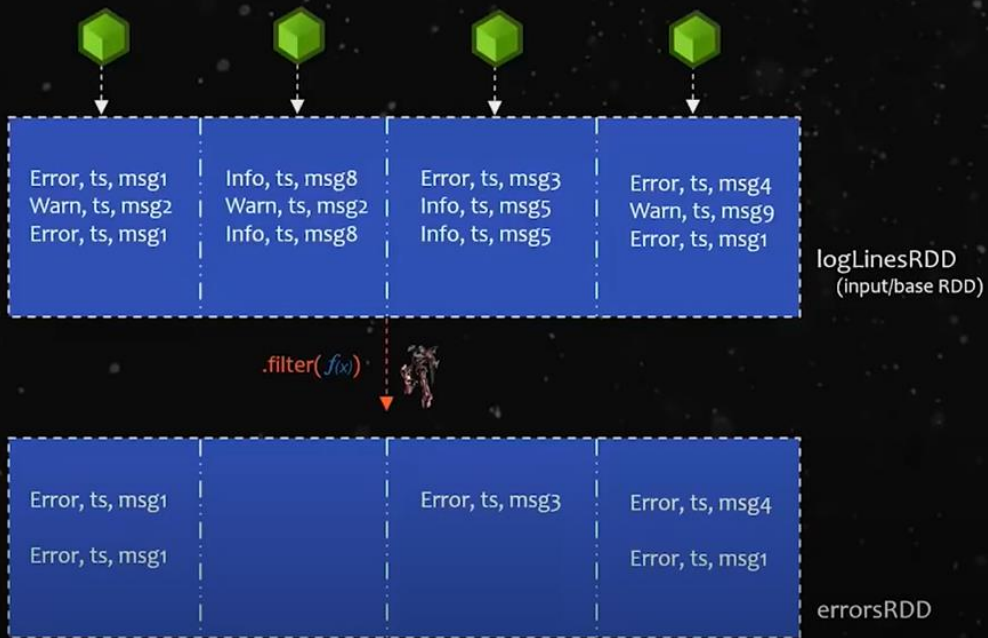


Example: performance

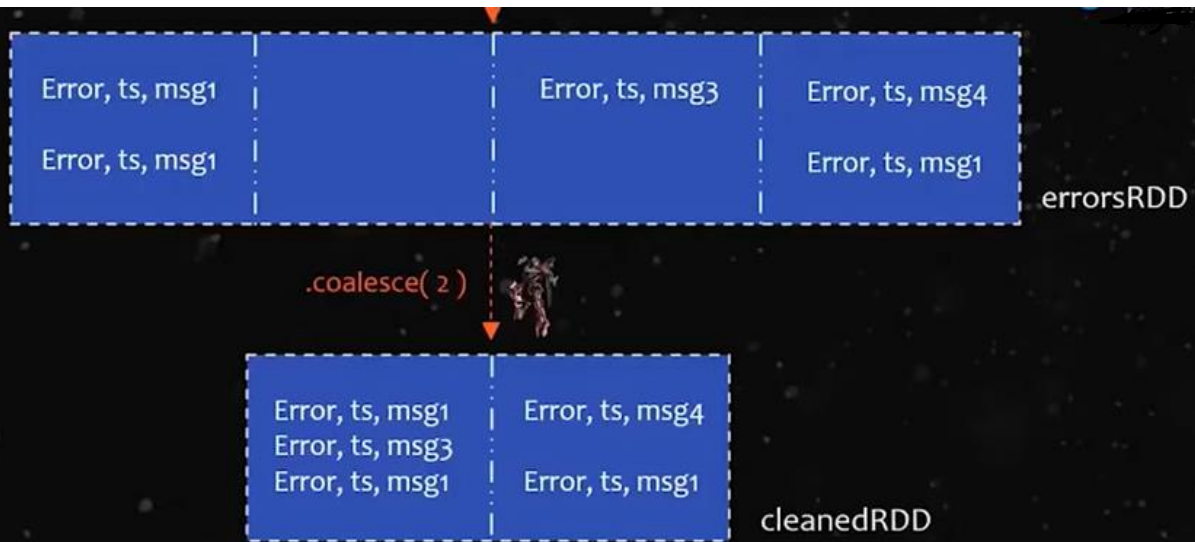
# Spark Operations

<b>Transformations</b> (create a new RDD)	map filter sample groupByKey reduceByKey sortByKey intersection	flatMap union join cogroup cross mapValues reduceByKey
<b>Actions</b> (return results to driver program)	collect Reduce Count takeSample take lookupKey	first take takeOrdered countByKey save foreach





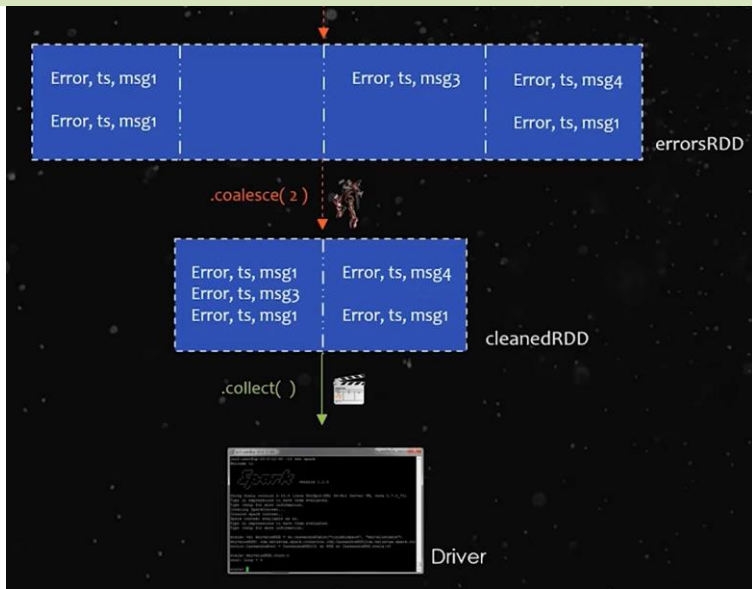
# Transformation




What is an action?

- The final stage of the workflow
- Triggers the execution of the DAG
- Returns the results to the driver
- Or writes the data to HDFS or to a file

# Spark Actions



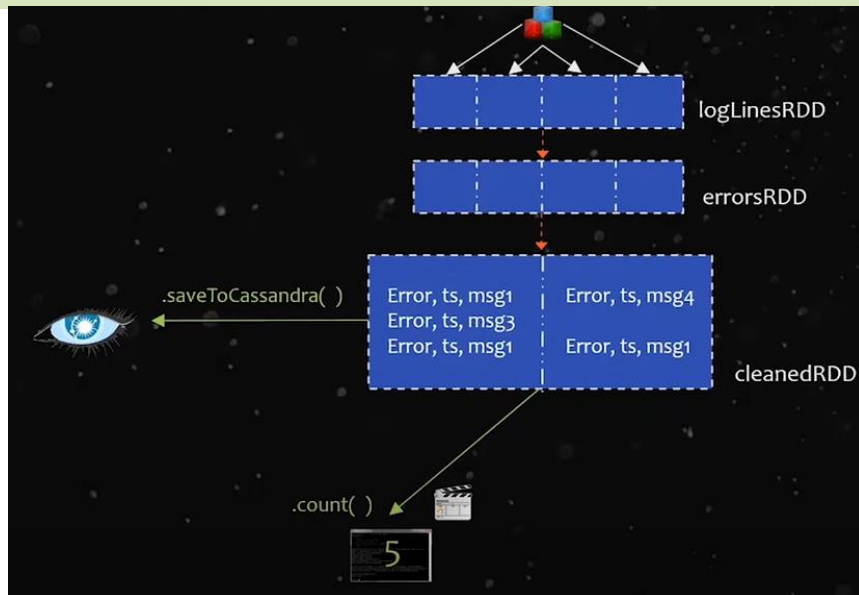
# Directed Acyclic Graphs (DAG)



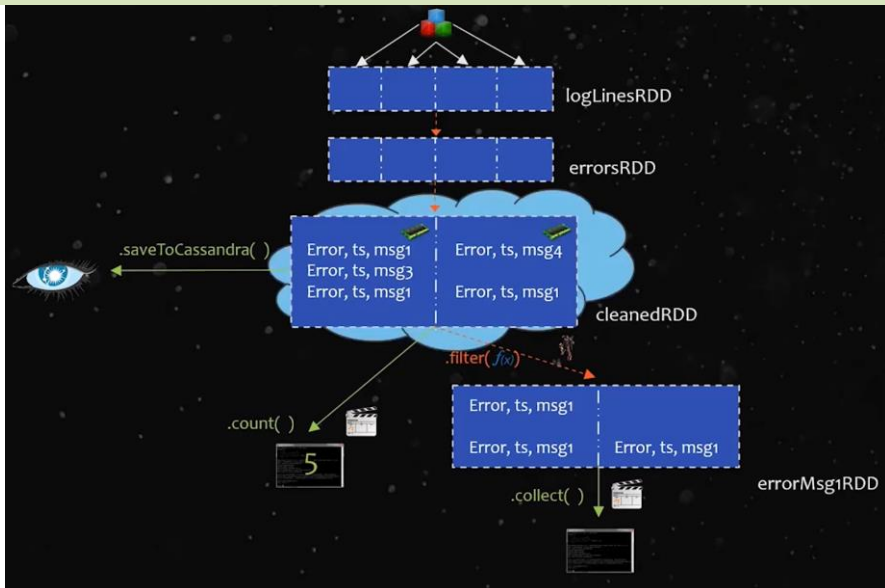
DAGs track dependencies (also known as Lineage )

- nodes are RDDs
- arrows are Transformations


# Spark DAG



# Spark DAG

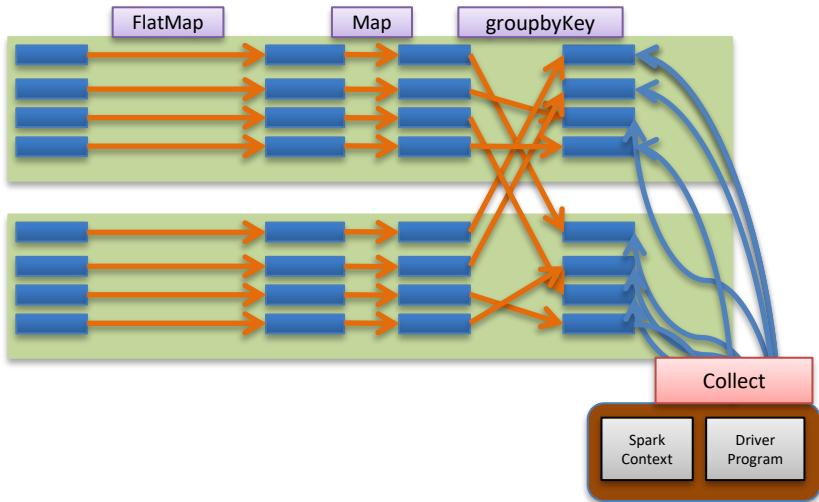


# Narrow Vs. Wide transformation





# Spark Workflow



# Python RDD API Examples

## Word count

```
text_file = sc.textFile("hdfs://usr/godil/text/book.txt")
counts = text_file.flatMap(lambda line: line.split(" ")) \
    .map(lambda word: (word, 1)) \
    .reduceByKey(lambda a, b: a + b)
counts.saveAsTextFile("hdfs://usr/godil/output/wordCount.txt")
```

## Logistic Regression

```
# Every record of this DataFrame contains the label and
# features represented by a vector.
df = sqlContext.createDataFrame(data, ["label", "features"])
# Set parameters for the algorithm.
# Here, we limit the number of iterations to 10.
lr = LogisticRegression(maxIter=10)
# Fit the model to the data.
model = lr.fit(df)
# Given a dataset, predict each point's label, and show the results.
model.transform(df).show()
```

Examples from <http://spark.apache.org/>

# RDD Persistence and Removal

## RDD Persistence

`RDD.persist()`

Storage level:

`MEMORY_ONLY, MEMORY_AND_DISK, MEMORY_ONLY_SER, DISK_ONLY, .....`

## RDD Removal

`RDD.unpersist()`

# Broadcast Variables and Accumulators (Shared Variables )

- Broadcast variables allow the programmer to keep a read-only variable cached on each node, rather than sending a copy of it with tasks

```
>broadcastV1 = sc.broadcast([1, 2, 3,4,5,6])
```

```
>broadcastV1.value
```

```
[1,2,3,4,5,6]
```

- Accumulators are variables that are only “added” to through an associative operation and can be efficiently supported in parallel

```
accum = sc.accumulator(0)
```

```
accum.add(x)
```

```
accum.value
```

# Spark's Main Use Cases

- Streaming Data
- Machine Learning
- Interactive Analysis
- Data Warehousing
- Batch Processing
- Exploratory Data Analysis
- Graph Data Analysis
- Spatial (GIS) Data Analysis
- And many more

# Spark in the Real World (I)

Uber – the online taxi company gathers terabytes of event data from its mobile users every day.

- By using Kafka, Spark Streaming, and HDFS, to build a continuous ETL pipeline
- Convert raw unstructured event data into structured data as it is collected
- Uses it further for more complex analytics and optimization of operations

Pinterest – Uses a Spark ETL pipeline

- Leverages Spark Streaming to gain immediate insight into how users all over the world are engaging with Pins—in real time.
- Can make more relevant recommendations as people navigate the site
- Recommends related Pins
- Determine which products to buy, or destinations to visit

# Spark in the Real World (II)

Here are Few other Real World Use Cases:

Conviva – 4 million video feeds per month

- This streaming video company is second only to YouTube.

- Uses Spark to reduce customer churn by optimizing video streams and managing live video traffic

- Maintains a consistently smooth, high quality viewing experience.

Capital One – is using Spark and data science algorithms to understand customers in a better way.

- Developing next generation of financial products and services

- Find attributes and patterns of increased probability for fraud

Netflix – leveraging Spark for insights of user viewing habits and then recommends movies to them.

- User data is also used for content creation

# Spark: when **not** to use

Even though Spark is versatile, that doesn't mean Spark's in-memory capabilities are the best fit for all use cases:


- For many simple use cases Apache MapReduce and Hive might be a more appropriate choice
- Spark was not designed as a multi-user environment
- Spark users are required to know that memory they have is sufficient for a dataset
- Adding more users adds complications, since the users will have to coordinate memory usage to run code



# HPC and Big Data Convergence

- Clouds and supercomputers are collections of computers networked together in a datacenter
- Clouds have different networking, I/O, CPU and cost trade-offs than supercomputers
- Cloud workloads are data oriented vs. computation oriented and are less closely coupled than supercomputers
- Principles of parallel computing same on both
- Apache Hadoop and Spark vs. Open MPI

# HPC and Big Data K-Means example



MPI definitely outpaces Hadoop, but can be boosted using a hybrid approach of other technologies that blend HPC and big data, including Spark and HARP. Dr. Geoffrey Fox, Indiana University. (<http://arxiv.org/pdf/1403.1528.pdf>)

# PGAS Vs MPI vs openMP

	Thread Count	Memory Count	Nonlocal Access
Traditional	1	1	N/A
OpenMP	Either 1 or p	1	N/A
MPI	p	p	No. Message required.
C+CUDA	1+p	2 (Host/device)	No. DMA required.
UPC, CAF, pMatlab	p	p	Supported.
X10, Asynchronous PGAS	p	q	Supported.